


✔ Congratulations! You passed!

Next Item

- ✔ 1 / 1 point
1. Consider using this encoder-decoder model for machine translation.



This model is a “conditional language model” in the sense that the encoder portion (shown in green) is modeling the probability of the input sentence  $x$ .

- ☐ True
- ☒ False

Correct

- ✔ 1 / 1 point
2. In beam search, if you increase the beam width  $B$ , which of the following would you expect to be true? Check all that apply.

☒ Beam search will run more slowly.

Correct

☒ Beam search will use up more memory.

Correct

☒ Beam search will generally find better solutions (i.e. do a better job maximizing  $P(y \mid x)$ )

Correct

☐ Beam search will converge after fewer steps.

Un-selected is correct

- ✔ 1 / 1 point
3. In machine translation, if we carry out beam search without using sentence normalization, the algorithm will tend to output overly short translations.

- ☒ True
- ☐ False

Correct

- ✔ 1 / 1 point
4. Suppose you are building a speech recognition system, which uses an RNN model to map from audio clip  $x$  to a text transcript  $y$ . Your algorithm uses beam search to try to find the value of  $y$  that maximizes  $P(y \mid x)$ .

On a dev set example, given an input audio clip, your algorithm outputs the transcript  $\hat{y}$  = “I’m building an A Eye system in Silly con Valley”, whereas a human gives a much superior transcript  $y^*$  = “I’m building an AI system in Silicon Valley”.

According to your model,

$$P(\hat{y} \mid x) = 1.09 \cdot 10^{-7}$$

$$P(y^* \mid x) = 7.21 \cdot 10^{-8}$$

Would you expect increasing the beam width  $B$  to help correct this example?

☒ No, because  $P(y^* \mid x) \leq P(\hat{y} \mid x)$  indicates the error should be attributed to the RNN rather than to the search algorithm.

Correct

☐ No, because  $P(y^* \mid x) \leq P(\hat{y} \mid x)$  indicates the error should be attributed to the search algorithm rather than to the RNN.

☐ Yes, because  $P(y^* \mid x) \leq P(\hat{y} \mid x)$  indicates the error should be attributed to the RNN rather than to the search algorithm.

☐ Yes, because  $P(y^* \mid x) \leq P(\hat{y} \mid x)$  indicates the error should be attributed to the search algorithm rather than to the RNN.


- ✔ 1 / 1 point
5. Continuing the example from Q4, suppose you work on your algorithm for a few more weeks, and now find that for the vast majority of examples on which your algorithm makes a mistake,  $P(y^* \mid x) > P(\hat{y} \mid x)$ . This suggest you should focus your attention on improving the search algorithm.

☒ True.

Correct

☐ False.

- ✔ 1 / 1 point
6. Consider the attention model for machine translation.



Further, here is the formula for  $\alpha^{<t,t'>}$ .

$$\alpha^{<t,t'>} = \frac{\exp(e^{<t,t'>})}{\sum_{t'=1}^{T_x} \exp(e^{<t,t'>})}$$

Which of the following statements about  $\alpha^{<t,t'>}$  are true? Check all that apply.

☒ We expect  $\alpha^{<t,t'>}$  to be generally larger for values of  $a^{<t'>}$  that are highly relevant to the value the network should output for  $y^{<t'>}$ . (Note the indices in the superscripts.)

Correct

☐ We expect  $\alpha^{<t,t'>}$  to be generally larger for values of  $a^{<t'>}$  that are highly relevant to the value the network should output for  $y^{<t'>}$ . (Note the indices in the superscripts.)

Un-selected is correct

☐  $\sum_t \alpha^{<t,t'>} = 1$  (Note the summation is over  $t$ .)

Un-selected is correct

☒  $\sum_{t'} \alpha^{<t,t'>} = 1$  (Note the summation is over  $t'$ .)

Correct

- ✔ 1 / 1 point
7. The network learns where to “pay attention” by learning the values  $e^{<t,t'>}$ , which are computed using a small neural network:

We can’t replace  $s^{<t-1>}$  with  $s^{<t>}$  as an input to this neural network. This is because  $s^{<t>}$  depends on  $\alpha^{<t,t'>}$  which in turn depends on  $e^{<t,t'>}$ ; so at the time we need to evaluate this network, we haven’t computed  $s^{<t>}$  yet.

- ☒ True
- ☐ False

Correct

- ✔ 1 / 1 point
8. Compared to the encoder-decoder model shown in Question 1 of this quiz (which does not use an attention mechanism), we expect the attention model to have the greatest advantage when:

☒ The input sequence length  $T_x$  is large.

Correct

☐ The input sequence length  $T_x$  is small.

- ✔ 1 / 1 point
9. Under the CTC model, identical repeated characters not separated by the “blank” character ( ) are collapsed. Under the CTC model, what does the following string collapse to?

\_c\_oo\_o\_kk\_\_b\_oooooo\_oo\_\_kkk

- ☐ cokbok
- ☒ cookbook

Correct

☐ cook book

☐ coookkboooooookkk

- ✔ 1 / 1 point
10. In trigger word detection,  $x^{<t>}$  is:

☒ Features of the audio (such as spectrogram features) at time  $t$ .

Correct

☐ The  $t$ -th input word, represented as either a one-hot vector or a word embedding.

☐ Whether the trigger word is being said at time  $t$ .

☐ Whether someone has just finished saying the trigger word at time  $t$ .